

```

import pandas as pd
from pandas.plotting import scatter_matrix
import numpy as np
import matplotlib.pyplot as plt
import os
from imblearn.over_sampling import ADASYN
from collections import Counter
import seaborn as sn
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.naive_bayes import BernoulliNB
from sklearn import metrics

```

```

pd.set_option('display.max_row',None)
pd.set_option('display.max_column',None)

```

```

def plot_confusion_matrix(cm, classes, title, cmap):
    "function for plotting confusion matrix"
    plt.clf()
    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    classnames = classes
    plt.title(title)
    plt.ylabel('True label')
    plt.xlabel('Predicted label')
    tick_marks = np.arange(len(classnames))
    plt.xticks(tick_marks, classnames, rotation=45)
    plt.yticks(tick_marks, classnames)
    s = [['TN', 'FP'], ['FN', 'TP']]
    for i in range(2):
        for j in range(2):
            plt.text(j, i, str(s[i][j])+" = "+str(cm[i][j]))
    plt.show()

```

```

def plot_roc_auc(model_list, X_test, y_test):
    "a function to plot roc_auc"
    fig, ax = plt.subplots(figsize=(8, 6))
    for model_name, model in model_list:
        y_score = model.predict_proba(X_test)[:, 1]
        fpr, tpr, _ = metrics.roc_curve(y_test, y_score)
        roc_auc = metrics.auc(fpr, tpr)
        plt.plot(fpr, tpr, lw=2, label=model_name + ' (area = %0.2f)' % roc_auc)
    plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
    plt.xlim([0.0, 1.0])
    plt.ylim([0.0, 1.05])
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title('ROC Curve')
    plt.legend(loc="lower right")
    plt.show()

```

```

import seaborn as sns

# Set up plotting parameters
%matplotlib inline
custom_style = "dark"
custom_palette = "colorblind"

```

```

# Set seaborn style and palette
sns.set_style(custom_style)
sns.set_palette(custom_palette)

```

```

df = pd.read_csv("creditcard.csv")

```

```

df.head(5)

```

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	V10	V11	V12	V13
0	0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363787	0.090794	-0.551600	-0.617801	-0.99131
1	0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425	-0.166974	1.612727	1.065235	0.48901
2	1	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654	0.207643	0.624501	0.066084	0.71721

```
# for determining the number of records in the dataset
print('The dataset contains {} rows and {} columns.'.format(df.shape[0], df.shape[1]))
```

The dataset contains 23858 rows and 31 columns.

```
# check for missing values and data types of the columns
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 23858 entries, 0 to 23857
Data columns (total 31 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Time        23858 non-null   int64
1   V1          23858 non-null   float64
2   V2          23858 non-null   float64
3   V3          23858 non-null   float64
4   V4          23858 non-null   float64
5   V5          23858 non-null   float64
6   V6          23858 non-null   float64
7   V7          23858 non-null   float64
8   V8          23858 non-null   float64
9   V9          23858 non-null   float64
10  V10         23858 non-null   float64
11  V11         23858 non-null   float64
12  V12         23858 non-null   float64
13  V13         23858 non-null   float64
14  V14         23858 non-null   float64
15  V15         23858 non-null   float64
16  V16         23858 non-null   float64
17  V17         23858 non-null   float64
18  V18         23858 non-null   float64
19  V19         23858 non-null   float64
20  V20         23858 non-null   float64
21  V21         23858 non-null   float64
22  V22         23857 non-null   float64
23  V23         23857 non-null   float64
24  V24         23857 non-null   float64
25  V25         23857 non-null   float64
26  V26         23857 non-null   float64
27  V27         23857 non-null   float64
28  V28         23857 non-null   float64
29  Amount      23857 non-null   float64
30  Class       23857 non-null   float64
dtypes: float64(30), int64(1)
memory usage: 5.6 MB
```

```
print('count of Normal transactions: ', df['Class'].value_counts().values[0])
print('count of Fraudulent transactions: ', df['Class'].value_counts().values[1])
```

```
count of Normal transactions: 23769
count of Fraudulent transactions: 88
```

```
# feature data (predictors)
features = df.iloc[:, :-1]
# label class
labels = df['Class']
```

```
scaler = StandardScaler()
scaled_features = scaler.fit_transform(features)
```

```
X_train, X_test, y_train, y_test = train_test_split(scaled_features, labels, test_size=0.30, random_state=42)
```

```
adasyn_sampler = ADASYN(random_state=42)
print('Original dataset shape {}'.format(Counter(y_train)))
X_resampled, y_resampled = adasyn_sampler.fit_resample(X_train, y_train)
print('Resampled dataset shape {}'.format(Counter(y_resampled)))
```

```
Original dataset shape Counter({0.0: 16636, 1.0: 64})
Resampled dataset shape Counter({0.0: 16636, 1.0: 16635})
```

```
X_train_resampled, y_train_resampled = X_resampled, y_resampled
```

```
# Train LogisticRegression Model
```

```
logistic_regression_classifier = LogisticRegression()
logistic_regression_classifier.fit(X_train_resampled, y_train_resampled)
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

```
Increase the number of iterations (max_iter) or scale the data as shown in:
```

```
https://scikit-learn.org/stable/modules/preprocessing.html
```

```
Please also refer to the documentation for alternative solver options:
```

```
https://scikit-learn.org/stable/modules/linear\_model.html#logistic-regression
```

```
n_iter_i = _check_optimize_result(
```

```
  ▾ LogisticRegression
```

```
LogisticRegression())
```

```
# Train Decision Tree Model
```

```
random_forest_classifier = RandomForestClassifier(random_state=0)
random_forest_classifier.fit(X_train, y_train)
```

```
  ▾ RandomForestClassifier
```

```
RandomForestClassifier(random_state=0)
```

```
# Train Bernoulli Naive Baye Model
```

```
bernoulli_nb_classifier = BernoulliNB()
bernoulli_nb_classifier.fit(X_train, y_train)
```

```
  ▾ BernoulliNB
```

```
BernoulliNB())
```

```
modlist = [('RandomForest Classifier', random_forest_classifier), ('LogisticRegression', logistic_regression_classifier),
            ('Naive Baye Classifier', bernoulli_nb_classifier)]
```

```
models = [j for j in modlist]
```

```
print()
```

```
print('===== Model Evaluation Results =====' "\n")
```

```
for model_name, model in models:
```

```
    scores = cross_val_score(model, X_train, y_train, cv=10)
```

```
    accuracy = metrics.accuracy_score(y_train, model.predict(X_train))
```

```
    confusion_matrix = metrics.confusion_matrix(y_train, model.predict(X_train))
```

```
    classification = metrics.classification_report(y_train, model.predict(X_train))
```

```
    print(f'==== {model_name} ====')
```

```
    print()
```

```
    print("Cross Validation Mean Score: ", '{:%'.format(np.round(scores.mean(), 3) * 100))
```

```
    print()
```

```
    print("Model Accuracy: ", '{:%'.format(np.round(accuracy, 3) * 100))
```

```
    print()
```

```
    print("Confusion Matrix:" "\n", confusion_matrix)
```

```
    print()
```

```
    print("Classification Report:" "\n", classification)
```

```
    print()
```

```
===== Model Evaluation Results =====
```

```
==== RandomForest Classifier =====
```

```
Cross Validation Mean Score: 99.9%
```

```
Model Accuracy: 100.0%
```

```
Confusion Matrix:
```

```
[[16636  0]
 [  0  64]]
```

```

Classification Report:
              precision    recall  f1-score   support

     0.0         1.00      1.00      1.00     16636
     1.0         1.00      1.00      1.00         64

 accuracy          1.00      1.00      1.00     16700
 macro avg          1.00      1.00      1.00     16700
 weighted avg          1.00      1.00      1.00     16700

```

```

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1)
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1)
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1)
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1)
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1)

```

```
df.describe()
```

	Time	V1	V2	V3	V4	V5
<b>count</b>	23858.000000	23858.000000	23858.000000	23858.000000	23858.000000	23858.000000
<b>mean</b>	18213.370609	-0.239141	0.198892	0.727022	0.248619	-0.188428
<b>std</b>	11377.032190	1.894219	1.533073	1.724887	1.440938	1.439894
<b>min</b>	0.000000	-30.552380	-40.978852	-31.103685	-5.172595	-42.147898
<b>25%</b>	6624.750000	-0.959528	-0.376134	0.287941	-0.658457	-0.767634
<b>50%</b>	20564.000000	-0.288644	0.192491	0.874426	0.216440	-0.218348
<b>75%</b>	29010.250000	1.164867	0.843146	1.505467	1.122367	0.325281
<b>max</b>	32954.000000	1.960497	16.713389	4.101716	11.927512	34.099309

